

Homework 9


Table of Contents

- [✓] ADSI Problem 6.1: Matched filters..... 1
 - 1) Plot a realisation of the signal with the added noise..... 2
 - 2) Plot the power density spectrum of the noise and spectrum of the signal..... 3
 - 3) Calculate the matched filter..... 5
 - 4) Does the frequency response of the matched filter make sense?..... 6
 - 5) How efficient is the matched filter?..... 7
 - 6) What happens to the optimum SNR when the noise is twice as powerful?..... 8
 - 7) Compute SNR using a 6-tap moving average filter..... 9
- Problem 14.14: M..... 10
 - [♦] 1) Determine the impulse response of the matched filter and the output SNR..... 11
 - [♦] 1.1) Explain why, in the absence of noise, the output is the ACRS of the desired signal..... 11
 - 2) Determine the matched filter of a short cosine signal..... 11
 - 3) Determine the matched filter of a long cosine signal..... 14
 - 4) Which signal can be detected more easily by visual inspection?..... 15
- Problem 14.35: Matched Filter and white noise..... 16
- Problem 14.37:..... 16
- Exam 2013, Problem 3..... 17
- Functions..... 17

```
figure('position', [0, 0, 800, 300])
```

[✓] ADSI Problem 6.1: Matched filters

Consider a deterministic signal that is only non-zero in the vicinity of $n=0$ as shown below and zero for all other values of n .



The signal is disturbed by additive noise from an AR(3) process given by

$$v(n) = -0.5v(n-1) - 0.5v(n-2) - 0.25v(n-3) + w(n)$$

where $w(n) \sim WGN(0, 4)$.

The goal of this problem is to construct a matched filter that can help us distinguish between the presence or absence of the signal in the noise.

```
clear variables;
```

1) Plot a realisation of the signal with the added noise

Create a realization of the noise and add the signal somewhere in the noise. Plot the result and comment on whether the signal is detectable.

```
D = 50; % The place to embed signal into the noise

% Generate the deterministic signal s[n]
s = [1, 4, 1, -1, -3, -2]';


% Generate white noise w[n] with zero mean and variance 4
w_var = 4;
w = sqrt(w_var) * randn(10000, 1);

% Generate noise from an AR(3) process
b = 1;
a = [1, 0.5, 0.5, 0.25];
v = filter(b, a, w);

% Extract 100 samples to remove the transient effect
% at the beginning and at the end of the filtered signal
x = v(500:600);

% Embed the signal s[n] in v[n] from D
s_in_x_start = D;
s_in_x_end = D + length(s) - 1;
x(s_in_x_start:s_in_x_end) = x(s_in_x_start:s_in_x_end) + s;

plot(1:length(x), x, 'c', ...
     s_in_x_start:s_in_x_end, s, 'r') % the signal is completely invisible
legend('s[n]+v[n]', 's[n]')
xlabel('n')
xlim([1,100])
ylim([-7, 9])
```



The signal is cannot be distinguished from the noise.

2) Plot the power density spectrum of the noise and spectrum of the signal

Plot the power density spectrum of the noise and spectrum of the signal, see Eq. (14.91).

The noise is generated by an AR(3) process plus some white Gaussian noise:

$$v(n) = -0.5v(n-1) - 0.5v(n-2) - 0.25v(n-3) + w(n)$$

where $w(n) \sim WGN(0, 4)$.

We can determine the Power Density Spectrum of an ARMA(p,q) process is given by

$$S_{yy}(\omega) = \sigma_x^2 |H(e^{j\omega})|^2 = \sigma_x^2 \left| \frac{\sum_{k=0}^q b_k e^{-j\omega k}}{1 + \sum_{k=1}^p a_k e^{-j\omega k}} \right|^2. \quad (13.133)$$

The power spectrum of an AR(p) process is given by:

$$S_{yy}(\omega) = \sigma_x^2 \left| \frac{1}{1 + \sum_{k=1}^p a_k e^{-j\omega k}} \right|^2$$

For this problem, we have:

$$S_{vv}(\omega) = 4 \left| \frac{1}{1 + 0.5e^{-j\omega} + 0.5e^{-j2\omega} + 0.25e^{-j3\omega}} \right|^2$$


The algorithm is as follows:

1. Use the coefficients $\{a_1, a_2, \dots, a_p\}$ for the AR(p) model,
2. Compute the transfer function for the AR(p) by computing the sum and finding its reciprocal
3. Compute the conjugate of the transfer function: $|H(e^{j\omega})|^2$

4. Multiply it with the variance σ_x^2


The algorithm is implemented in the functions `ar2psd()` function:

```
N = 256;  
  
a = [0.5, 0.5, 0.25]; % The coefficients of the AR(3) model  
w_var = 4; % The variance of white noise  
[S_vv, w] = ar2psd(a, w_var, N); % Compute the PDS of AR(3) model  
  
plot(w/pi, pow2db(S_vv))  
xlabel('\omega/\pi')  
ylabel('S_{vv}^{AR(3)}(\omega) in db')
```



We can plot the spectrum of the signal $s(n)$ using the `freqz()` function:

```
[H,w2] = freqz(s, 1);  
plot(w2/pi, 20*log10(abs(H)))  
xlabel('\omega/\pi')  
ylabel('Magnitude Response (dB)')
```



3) Calculate the matched filter

Calculate the matched filter, the frequency response of the matched filter and the optimum signal to noise ratio. You can use `xcorr` on the noise realization to find \mathbf{R}_v .

The impulse response of the matched filter is given by:

$$\mathbf{h} = \kappa \mathbf{R}_v^{-1} \mathbf{s}. \quad (14.97)$$

where k is the normalisation factor. Although the maximum SNR can be obtained by any choice of constant κ , we choose the constant by requiring that:

- (a) $\mathbf{h}^T \mathbf{s} = 1$, which yields $\kappa = 1/\mathbf{s}^T \mathbf{R}_v^{-1} \mathbf{s}$
- (b) $E(v_0^2[n]) = \mathbf{h}^T \mathbf{R}_v \mathbf{h} = 1$, which yields $\kappa = 1/\sqrt{\mathbf{s}^T \mathbf{R}_v^{-1} \mathbf{s}}$.


```
p = numel(s); % Signal length

% The autocorrelation matrix must be p x p since
% its inverse is multiplied by a p-tap signal s(n)
[r_vv, ~] = xcorr(v, p-1, 'biased');
R_vv = toeplitz(r_vv(p:end));

% Compute normalisation factor
k = 1/sqrt(s'*(R_vv\s)); % Same as 1/sqrt(s'*inv(R_vv)*s)

% Compute the filter
h = k*(R_vv\s); % Same as k*inv(R_vv)*s

[H_s, w_s] = freqz(s, 1);
[H_h, w_h] = freqz(h, 1);
plot(w_s/pi, 20*log10(abs(H_s)), w_h/pi, 20*log10(abs(H_h)))
legend('Signal', 'Matched Filter')
xlabel('\omega/\pi')
ylabel('Magnitude Response (dB)')
grid on;
```



The maximum possible value of the output SNR is given by:

$$\text{SNR}_0 = a^2 \tilde{s}^T \tilde{s} = a^2 s^T R_v^{-1} s. \quad (14.98)$$

Since the attenuation factor a is not given in this problem, we assume that it is 1:


```
a = 1;
SNR = a^2 * s' * (R_vv\s)
```

```
SNR = 11.5754
```

4) Does the frequency response of the matched filter make sense?


Does the frequency response of the matched filter make sense seen in relation with the spectra from question 2?

In question 2) we found that the AR(3) noise model has peak power at 0.5.



Looking at the frequency response of the signal, we observe that it peaks around 0.25 and dips around 20 dB at around 0.6.

The frequency response of the matched filter shows that it is minimising the noise in the region at around 0.55. The noise power peaks at around 0.5



5) How efficient is the matched filter?

Plot the signal before and after the matched filter as well as the square of the output of the matched filter on the same graph and comment on the efficiency of the matched filter in our quest to detect the presence of the signal.


```
% x(n) is the 100 samples extracted from the noise signal v(n)
% y(n) is the result of feeding x(n) to the matched filter
y = filter(h, 1, x);
y_squared = y.^2;

n = 1:length(x);


n2 = s_in_x_start:s_in_x_end;
xs = s;

plot(n,x, n,y, n,y_squared, n2,xs,'*-')
xlim([0, numel(n)])
ylim([-max(y_squared), max(y_squared)])
legend({'x[n] Measured signal', 'y[n] Matched filter output', 'y^2[n]', 's[n] Original signal'},
'Location','southwest')

grid on;
```



It's close to impossible to find the original signal $s[n]$ in the measured signal $x[n]$. It is also difficult to detect $s[n]$ after filtering $x[n]$ with the matched filter i.e., $y[n]$. It is easier to detect the original signal in the square of the output of the matched filter $y^2[n]$. Running the code multiple times i.e., creating different realisations of the noise, we may see many false positives. The figure below shows an example of this. The highest response is measured at $n = 6$ but we know that the signal is embedded in $n = 50$.



6) What happens to the optimum SNR when the noise is twice as powerful?

Is the optimum signal to noise ratio halved if the AR(3) is instead driven by twice as powerful white noise, i.e. $w(n) \sim \text{WGN}(0,8)$? Why or why not?

We will compute the SNR again but this time with $\sigma_w^2 = 8$:

$$\text{SNR}_0 = a^2 \tilde{s}^T \tilde{s} = a^2 s^T R_v^{-1} s. \quad (14.98)$$

```
% Generate white noise w[n] with zero mean and variance 8
w2_var = 8;
w2 = sqrt(w2_var) * randn(10000, 1);

% Put the white noise through an AR(3) process
b = 1;
a = [1, 0.5, 0.5, 0.25];
v2 = filter(b, a, w2);

% Extract 100 samples to remove the transient effect
% at the beginning and at the end of the filtered signal
x2 = v2(500:600);

p = numel(s); % Signal length

% The autocorrelation matrix must be MxM since
% its inverse is multiplied by a M-tap signal s(n)
[r_vv2, lags] = xcorr(x2, p-1, 'biased');
R_vv2 = toeplitz(r_vv2(p:end));

a = 1;
```


$$\text{SNR2} = a^2 * s' * (R_{vv2} \backslash s)$$

$$\text{SNR2} = 6.6825$$

$$\text{SNR2}/\text{SNR}$$

$$\text{ans} = 0.5773$$

The optimal signal-to-noise-ratio is almost halved when the white noise is doubled. The reduction of SNR is not always decreased by a factor of though. Although the noise power is doubled, the shape of the optimum filter will change because it depends on the measured signal:

$$\mathbf{h} = \kappa \mathbf{R}_v^{-1} \mathbf{s}. \quad (14.97)$$

7) Compute SNR using a 6-tap moving average filter

Assume that instead of the matched filter a simple 6-tap moving average filter is used instead. That is

$$\mathbf{h} = \left[\frac{1}{6} \quad \frac{1}{6} \quad \frac{1}{6} \quad \frac{1}{6} \quad \frac{1}{6} \quad \frac{1}{6} \right]^T$$

Calculate the signal to noise ratio when the moving average filter is used.

We can compute the SNR using the following formula:

$$\text{SNR}_0 = \frac{s_o^2[n_0]}{E(v_o^2[n_0])} = a^2 \frac{(\mathbf{h}^T \mathbf{s})^2}{\mathbf{h}^T \mathbf{R}_v \mathbf{h}}. \quad (14.94)$$

```
% h_avg = [1, 0, 1, 1, 1, 1]'./5;
h_avg = [1, 1, 1, 1, 1, 1]'./6;
a = 1;
SNR_avg = a^2 * (h_avg'*s)^2 / (h_avg'*R_vv*h_avg)
```

$$\text{SNR}_{\text{avg}} = 0$$

The SNR of the moving average filter is zero (or a number very close to zero). The moving average filter gives us the average of a signal. Since the original signal has zero mean, the average is zero. This is not surprising. If we change the filter, then the SNR is no longer zero.

```
h4 = [1, 0, 1, 1, 1, 1]'./5;
a = 1;
SNR4 = a^2 * (h4'*s)^2 / (h4'*R_vv*h4)
```

$$\text{SNR4} = 1.4075$$


Problem 14.14: Matched filters, compare short vs long signals

```
clear variables;
```


In this problem we discuss in detail the matched filtering problem illustrated in Figure 14.18.

Figure 14.18 shows the operation of a matched filter in white noise.

The two signals have different lengths ($p = 10$ and $p = 100$, respectively) but the same energy.



Notice that the peak response of the matched filter occurs at $n_0 = p - 1$ because there is no time delay i.e. $D = 0$. Normally, the output SNR is maximised at time $n_0 = p + D - 1$. In practice, finding the time delay D depends on the particular application.



The input to the matched filter is given by

$$x[n] = s_I[n] + v[n] = a s[n - D] + v[n]$$

where

- a is the attenuation factor
- D is the round-trip delay
- $v[n]$ is the white Gaussian noise with zero mean and unit variance
- $s_I[n] = 0$ outside the interval $0 \leq n \leq p - 1$

[◆] 1) Determine the impulse response of the matched filter and the output SNR.

Determine the impulse response of the matched filter and the output SNR. Explain why, in the absence of noise, the output is the ACRS of the desired signal.

[◆] To compute the matched filter and the output SNR, we need to know something about the signal $s[n]$ and the noise.

[◆] 1.1) Explain why, in the absence of noise, the output is the ACRS of the desired signal.

In case of white noise, the impulse response of the matched filter $h[n]$ is equal to the signal that we want to detect $s[n]$. As such the output of the matched filter is the same as autocorrelation:

2) Determine the matched filter of a short cosine signal

Suppose that $p = 10$ and $s_i[n] = \cos\left(2\pi \frac{n}{10}\right)$. Generate $N = 200$ samples of the noisy signal $s[n]$ and process it through the matched filter designed for $s_i[n]$. Plot the desired, input, and filtered signals and determine when the matched filter output is output.

The impulse response of the matched filter is given by:

$$h = \kappa R_v^{-1} s. \quad (14.97)$$

where k is the normalisation factor and R_v is the autocorrelation matrix of the noise.

The output SNR can be computed using the formula:

$$\text{SNR}_o = a^2 \tilde{s}^T \tilde{s} = a^2 s^T R_v^{-1} s. \quad (14.98)$$

The computation is coded in a MATLAB function named `matched_filter` (see the end of the document)

```
N = 200;  
n = (1:N)';
```

```

p = 10; % Signal length

% Generate the attenuated signal.
% Ensure that s[n] is zero when n is outside the interval [1, p]
s = zeros(N, 1);
inside_indices = (n >= 1 & n <= p);
s(inside_indices) = cos(2*pi*n(inside_indices)/10);

% Generate the zero-mean WGN with unit variance
v = randn(N, 1);


% Generate the measured signal x[n]
x = v + s;

% Compute the impulse response of the matched filter
[h, SNR] = matched_filter(s(1:p), v);

% Feed the measured signal to the matched filter
y = filter(h, 1, x);

figure('position', [0, 0, 800, 800])
subplot(3,1,1); plot(n,s); ylabel('Original signal s[n]')
subplot(3,1,2); plot(n,x); ylabel('Measured signal x[n]')
subplot(3,1,3); plot(n,y); ylabel('Filter output y[n]')
xlim([0 N])
grid on;


```



From the filtered output, it is difficult to determine where the original signal is placed.

Let us compare the impulse response of the filter to the signal:

```
figure('position', [0, 0, 800, 300])
n2 = 1:p;
stem(n2, s(n2))
hold on; stem(n2, h); hold off;
legend('s[n]', 'h[n]', 'Location', 'north')
```



3) Determine the matched filter of a long cosine signal

Repeat 2) for $p = 100$ and $s_i[n] = \frac{1}{\sqrt{10}} \cos\left(2\pi \frac{n}{10}\right)$.

```
p = 100; % Signal length

% Generate the attenuated signal.
% Ensure that s[n] is zero when n is outside the interval [1, p]
s = zeros(N, 1);
inside_indices = (n >= 1 & n <= p);
s(inside_indices) = (1/sqrt(10)) * cos(2*pi*n(inside_indices)/10);

% Generate the zero-mean WGN with unit variance
v = randn(N, 1);


% Generate the measured signal x[n]
x = v + s;

% Compute the impulse response of the matched filter
[h_100, SNR_100] = matched_filter(s(1:p), v);

% Feed the measured signal to the matched filter
y = filter(h_100, 1, x);

figure('position', [0, 0, 800, 800])
subplot(3,1,1); plot(n,s); ylabel('Original signal s[n]')
subplot(3,1,2); plot(n,x); ylabel('Measured signal x[n]')
subplot(3,1,3); plot(n,y); ylabel('Filter output y[n]')
```


```
xlim([0 N])  
grid on;
```



4) Which signal can be detected more easily by visual inspection?

Which signal can be detected more easily by visual inspection of the matched filter output? Is this justified by comparing the output SNR in each case?

For certain noise realisations, it is easier to detect the second (longer) signal by visual inspect than the first signal. However, running the code multiple times i.e., creating different realisations of the noise, we may see many false positives. The figure below shows an example of this. The highest response is measured at $n = 140$ but the maximum filter response should occur at $n = p + D - 1 = 100 + 0 - 1 = 99$.



Comparing the output SNR of the two filters, we observe that the SNR of the longer filter is usually better but this is not always the case. Different realisation of the noise process, yields different results:

SNR

SNR = 4.7238

SNR_100

SNR_100 = 5.5132

Problem 14.35: Matched Filter and white noise

In the book on page 861, the authors claim: "*The term matched filter was introduced because, in the case of white noise, the impulse response is "matched" to the shape of the signal being sought. In fact, the output of the matched filter is proportional to the correlation between the signal segment stored in the filter memory and the signal of interest.*"

This problem investigates this claim.

Consider the two 10-point signals $x_0[n]$ and $x_1[n]$ given below:

```
clear variables;
```

Problem 14.37:


```
clear variables;
```

Exam 2013, Problem 3

```
clear variables;
```

Functions

```
function [S, w] = ar2psd(a, v, N)
% AR2PSD Compute the Power Spectral Density from AR(p) coefficients
% [S, w] = ar2psd(a, v, N)
% a: AR(p) coefficients
% v: the variance
% N: number of points in the range [1, pi]
% S: the estimated power spectrum
% w: frequencies
    w = linspace(0, 1, N) * pi;

    % Compute the transfer function
    % Used Eq. (13.133) in the book
    H = ones(N, 1);
    for k=1:numel(a)
        H = H + a(k)*exp(-1j * w' * k);
    end
    H = 1./H;

    % Finally compute the PSD
    S = v * H.*conj(H);
end

function [h, SNR] = matched_filter(s, v, a)
% MATCHED_FILTER: Compute the impulse response of a matched filter and
%                  and the corresponding output SNR.
```

```

% [h, SNR] = matched_filter(s, v, a)
% s: the signal
% v: a realisation of the additive noise
% a: the attenuation factor (default=1)
% h: the impulse response of the matched filter
% SNR: the output SNR
    if nargin < 3
        a = 1;
    end

    p = numel(s); % Signal length

    % The autocorrelation matrix must be p x p since
    % its inverse is multiplied by a p-tap signal s(n)
    [r_vv, ~] = xcorr(v, p-1, 'biased');
    R_vv = toeplitz(r_vv(p:end));

    % The expression `R_vv^{-1} * s` is used multiple times,
    % so compute it once and reuse
    R_vv_inv_s = R_vv\s;

    % Compute normalisation factor
    k = 1/sqrt(s'*R_vv_inv_s);

    % Compute the filter
    h = k*R_vv_inv_s;

    % Compute the output SNR
    SNR = a^2 * s' * R_vv_inv_s;
end

```